

# Package ‘LinkOrgs’

December 23, 2025

**Title** LinkOrgs: Algorithms for Organizational Record Linkage

**Version** 0.01

**Author** Connor Jerzak [aut, cre] (<<https://orcid.org/0000-0003-1914-8905>>),  
Brian Libgober [aut] (<<https://orcid.org/0000-0001-9638-0228>>)

**Maintainer** Connor Jerzak <connor.jerzak@gmail.com>

**Description** An R package for organizational records using the algorithms of Jerzak & Libgober (2023+). The linkage is done based on organizational names and using half a billion open collaborated records on those names from LinkedIn users. It also contains functions implementing string matching performance metrics, as well as a fast, parallelized version of fuzzy string matching.

**Depends** R (>= 3.5.0)

**License** CC BY-NC-ND 4.0

**Encoding** UTF-8

**Imports** data.table,plyr,stringdist,parallel,stringr,dplyr,fastmatch,reticulate,foreach,doParallel,digest

**Suggests** testthat (>= 3.0.0), knitr, rmarkdown, text

**VignetteBuilder** knitr

**Config/testthat/edition** 3

**RoxygenNote** 7.3.3

**NeedsCompilation** no

## Contents

|                                  |    |
|----------------------------------|----|
| AssessMatchPerformance . . . . . | 2  |
| BuildBackend . . . . .           | 3  |
| dropboxURL2downloadURL . . . . . | 4  |
| GetCalibratedDistThres . . . . . | 5  |
| LinkOrgs . . . . .               | 6  |
| pDistMatch_discrete . . . . .    | 9  |
| pDistMatch_euclidean . . . . .   | 11 |
| pFuzzyMatch_discrete . . . . .   | 12 |
| pFuzzyMatch_euclidean . . . . .  | 14 |
| print2 . . . . .                 | 16 |
| url2dt . . . . .                 | 16 |

## Index

18

AssessMatchPerformance  
*AssessMatchPerformance*

## Description

Computes match performance metrics (true/false positive and true/false negative rates) by comparing predicted matches against a ground-truth dataset.

## Usage

```
AssessMatchPerformance(
  x,
  y,
  z,
  z_true,
  by,
  by.x = by,
  by.y = by,
  openBrowser = F
)
```

## Arguments

|             |   |
|-------------|---|
| x, y        | Data frames that were merged to produce z. Used to determine the universe of possible match pairs.            |
| z           | Data frame containing the predicted matches to be evaluated. Must contain columns specified by by.x and by.y. |
| z_true      | Data frame containing the ground-truth (reference) matches. Must contain columns specified by by.x and by.y.  |
| by          | Character string specifying the column name for merging when both data frames use the same column name.       |
| by.x        | Character string specifying the column name in x used for merging. Defaults to by if not specified.           |
| by.y        | Character string specifying the column name in y used for merging. Defaults to by if not specified.           |
| openBrowser | Logical; if TRUE, opens browser for debugging. Default is FALSE.  |

## Value

A named numeric vector with the following elements:

**TruePositives** Number of matches in z that are also in z\_true.

**FalsePositives** Number of matches in z that are not in z\_true.

**FalseNegatives** Number of matches in z\_true that are not in z.

**TrueNegatives** Number of non-matches correctly identified (total pairs minus TP, FP, and FN).

**MatchedDatasetSize** Number of rows in the predicted matches z.

## Examples

```
# Create synthetic data
x_ornames <- c("apple", "oracle", "enron inc.", "mcdonalds corporation")
y_ornames <- c("apple corp", "oracle inc", "enron", "mcdonalds co")
x <- data.frame("orgnames_x" = x_ornames)
y <- data.frame("orgnames_y" = y_ornames)
z <- data.frame("orgnames_x" = x_ornames[1:2], "orgnames_y" = y_ornames[1:2])
z_true <- data.frame("orgnames_x" = x_ornames, "orgnames_y" = y_ornames)

# Obtain match performance data
PerformanceMatrix <- AssessMatchPerformance(x = x,
                                              y = y,
                                              z = z,
                                              z_true = z_true,
                                              by.x = "orgnames_x",
                                              by.y = "orgnames_y")
print(PerformanceMatrix)
```

## Description

Creates and configures a conda environment with all necessary Python packages (JAX, TensorFlow, Optax, Equinox, JMP) for running the machine learning components of LinkOrgs.

## Usage

```
BuildBackend(conda_env = "LinkOrgs_env", conda = "auto", tryMetal = T)
```

## Arguments

|           |  |
|-----------|--|
| conda_env | Character string; name of the conda environment to create. Default is "LinkOrgs_env".  |
| conda     | Character string; path to a conda executable, or "auto" to let reticulate automatically find an appropriate conda binary. Default is "auto". |
| tryMetal  | Logical; if TRUE and running on Apple Silicon (arm64 macOS), attempts to install jax-metal for GPU acceleration. Default is TRUE.            |

## Details

This function requires an Internet connection to download packages. The conda environment will include:

- TensorFlow 2.15
- TensorFlow Probability 0.23
- JAX 0.4.26 and JAXlib 0.4.26
- Optax 0.2.2
- Equinox 0.11.4
- JMP 0.0.4

- NumPy 1.26.4
- jax-metal 0.1.0 (Apple Silicon only, if `tryMetal = TRUE`)

You can find available conda Python paths via: `system("which python")`

### Value

Invisibly returns NULL. Called for its side effect of creating and configuring the conda environment.  
Prints "Done building LinkOrgs backend!" upon successful completion.

### See Also

[LinkOrgs\(\)](#) for using the ML backend after setup.

### Examples

```
## Not run:
# Build with default settings
BuildBackend()

# Build with a specific conda path
BuildBackend(conda = "/opt/miniconda3/bin/conda")

# Build without attempting Metal support on macOS
BuildBackend(tryMetal = FALSE)

## End(Not run)
```

## dropboxURL2downloadURL

*Convert Dropbox Share URL to Direct Download URL*

### Description

Converts a Dropbox share link to a direct download URL by replacing the domain with `d1.dropboxusercontent.com`.

### Usage

`dropboxURL2downloadURL(url)`

### Arguments

|                  |   |
|------------------|---|
| <code>url</code> | Character string; a Dropbox share URL (e.g., " <code>https://www.dropbox.com/s/...</code> " or " <code>https://dropbox.com/s/...</code> "). |
|------------------|---|

### Details

Dropbox share links require modification to enable direct file downloads. This function replaces:

- `https://www.dropbox.com` with `https://d1.dropboxusercontent.com`
- `www.dropbox.com` with `d1.dropboxusercontent.com`
- `dropbox.com` with `d1.dropboxusercontent.com`

**Value**

Character string; the converted direct download URL. If the input is not a Dropbox URL, it is returned unchanged.

**See Also**

[url2dt\(\)](#) which uses this function internally.

**Examples**

```
# Convert a Dropbox share link
direct_url <- dropboxURL2downloadURL(
  "https://www.dropbox.com/s/abc123/myfile.csv?dl=0"
)
# Returns: "https://dl.dropboxusercontent.com/s/abc123/myfile.csv?dl=0"
```

**GetCalibratedDistThres**

*GetCalibratedDistThres*

**Description**

Calibrates a distance threshold based on a target average number of matches per alias. Samples pairwise distances from a subset of observations to determine the threshold that would yield approximately the desired number of matches.

**Usage**

```
GetCalibratedDistThres(
  x = NULL,
  y = NULL,
  by.x = NULL,
  by.y = NULL,
  AveMatchNumberPerAlias = 5L,
  qgram = 2L,
  DistanceMeasure = "jaccard",
  nCores = NULL,
  mode = "euclidean"
)
```

**Arguments**

|      |  |
|------|--|
| x    | Input data. For mode = "euclidean": an embedding matrix where rows are observations and columns are embedding dimensions. For mode = "discrete": a data frame containing the column specified by by.x. |
| y    | Input data. For mode = "euclidean": an embedding matrix where rows are observations and columns are embedding dimensions. For mode = "discrete": a data frame containing the column specified by by.y. |
| by.x | Column name in x to use for matching. Only used when mode = "discrete".  |
| by.y | Column name in y to use for matching. Only used when mode = "discrete".  |

|                        |   |
|------------------------|---|
| AveMatchNumberPerAlias | Target average number of matches per observation. Used to calibrate the distance threshold. Default is 5.   |
| qgram                  | The q-gram size for string distance calculation. Only used when mode = "discrete". Default is 2.  |
| DistanceMeasure        | The string distance measure to use. Only used when mode = "discrete". Options include "jaccard", "osa", "jw". See ?stringdist::stringdist for all options. Default is "jaccard".        |
| nCores                 | Number of CPU cores for parallel computation. Only used when mode = "discrete". Default is NULL (auto-detect).  |
| mode                   | Character string specifying the distance computation mode. Must be either "euclidean" (for embedding-based matching) or "discrete" (for string-based matching). Default is "euclidean". |

**Value**

A numeric value representing the calibrated distance threshold.

LinkOrgs

*LinkOrgs***Description**

Implements the organizational record linkage algorithms of Libgober and Jerzak (2023+) using half-a-billion open-collaborated records.

**Usage**

```
LinkOrgs(
  x = NULL,
  y = NULL,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  embedx = NULL,
  embedy = NULL,
  embedDistMetric = NULL,
  algorithm = "ml",
  conda_env = "LinkOrgs_env",
  conda_env_required = T,
  ReturnDiagnostics = F,
  ReturnProgress = T,
  ToLower = T,
  NormalizeSpaces = T,
  RemovePunctuation = T,
  MaxDist = NULL,
  MaxDist_network = NULL,
  AveMatchNumberPerAlias = 10,
  AveMatchNumberPerAlias_network = 2,
```

```

  DistanceMeasure = "jaccard",
  qgram = 2,
  RelThresNetwork = 1.5,
  ml_version = "v1",
  openBrowser = F,
  ExportEmbeddingsOnly = FALSE,
  ReturnDecomposition = FALSE,
  python_executable,
  nCores = NULL,
  deezyLoc = NULL
)

```

## Arguments

|                                 |  |
|---------------------------------|--|
| <code>x, y</code>               | Data frames to be merged.  |
| <code>by, by.x, by.y</code>     | Character vector(s) that specify the column names used for merging data frames <code>x</code> and <code>y</code> . The merging variables should be organizational names. See <code>?base::merge</code> for more details regarding syntax.  |
| <code>embedx, embedy</code>     | Optional pre-computed embedding matrices. If provided, these will be used instead of computing embeddings from names. Rows correspond to observations and columns to embedding dimensions.   |
| <code>embedDistMetric</code>    | Optional custom distance metric function for embedding-based matching.   |
| <code>algorithm</code>          | Character; specifies which algorithm should be used. Options are "fuzzy", "ml", "bipartite", "markov", and "transfer". Default is "ml", which uses a machine-learning approach with Transformer networks and up to 11 million parameters to predict match probabilities using half a billion open-collaborated records as training data. |
| <code>conda_env</code>          | Character string; specifies a conda environment where JAX and related packages have been installed (see <code>?LinkOrgs::BuildBackend</code> ). Used only when <code>algorithm='ml'</code> or <code>DistanceMeasure='ml'</code> . Default is "LinkOrgs_env".   |
| <code>conda_env_required</code> | Logical; specifies whether conda environment is required. Default is TRUE.   |
| <code>ReturnDiagnostics</code>  | Logical; if TRUE, various match-level diagnostics are returned in the merged data frame. Default is FALSE.   |
| <code>ReturnProgress</code>     | Logical; if TRUE, progress messages are printed during execution. Default is TRUE.   |
| <code>ToLower</code>            | Logical; if TRUE, converts names to lowercase before matching. Default is TRUE.  |
| <code>NormalizeSpaces</code>    | Logical; if TRUE, removes extra whitespace from names. Default is TRUE.  |
| <code>RemovePunctuation</code>  | Logical; if TRUE, removes punctuation from names. Default is TRUE.   |
| <code>MaxDist</code>            | Numeric; maximum allowed distance between two matched strings. If <code>AveMatchNumberPerAlias</code> is specified, it takes priority over this parameter.   |
| <code>MaxDist_network</code>    | Numeric; maximum allowed distance for network-based matching when using <code>algorithm = "bipartite"</code> or <code>"markov"</code> .  |

|                                |  |
|--------------------------------|--|
| AveMatchNumberPerAlias         | Numeric; target average number of matches per alias. Used to automatically calibrate MaxDist. Takes priority over MaxDist if both are specified. Default is 10.  |
| AveMatchNumberPerAlias_network | Numeric; target average number of matches per alias for network-based candidate selection. Default is 2.   |
| DistanceMeasure                | Character; algorithm for computing pairwise string distances. Options include "osa", "jaccard", "jw", or "ml" for embedding-based distance. See ?stringdist::stringdist for all string distance options. Default is "jaccard". |
| qgram                          | Integer; the q-gram size used in string distance measures. Default is 2.   |
| RelThresNetwork                | Numeric; relative threshold multiplier for network distances. Default is 1.5.  |
| ml_version                     | Character; specifies which version of the ML algorithm to use. Options are "v0" (9M parameters) through "v4". Default is "v1" (11M parameters).  |
| openBrowser                    | Logical; if TRUE, opens browser for debugging. Default is FALSE.   |
| ExportEmbeddingsOnly           | Logical; if TRUE with algorithm='ml' (or DistanceMeasure='ml'), returns only ML embeddings for x and/or y without matching, for offline linkage. Default is FALSE.   |
| ReturnDecomposition            | Logical; if TRUE, returns a list containing the merged data frame along with intermediate results. Default is FALSE.   |
| python_executable              | Path to Python executable. Usually not needed if conda_env is specified.   |
| nCores                         | Integer; number of CPU cores to use for parallel processing. Default is NULL (auto-detect based on data size).   |
| deozyLoc                       | Path to DeezyMatch installation (for algorithm = "deezymatch").  |

## Details

LinkOrgs automatically processes the name text for each dataset (specified by by or by.x and by.y). Text preprocessing includes:

- **Case normalization:** Set ToLower = FALSE to preserve case sensitivity.
- **Space normalization:** Set NormalizeSpaces = FALSE to preserve whitespace.
- **Punctuation removal:** Set RemovePunctuation = FALSE to preserve punctuation.

To use combined machine learning and network methods, set algorithm to "bipartite" or "markov", and DistanceMeasure to "ml".

## Value

If ExportEmbeddingsOnly = TRUE, returns a list with embedx and/or embedy data frames containing the input names and their embeddings. If ReturnDecomposition = TRUE, returns a list with z (merged data), z\_RawNames (raw name matches), and z\_Network (network matches). Otherwise, returns the merged data frame z.

## Examples

```
# Create synthetic data
x_ornames <- c("apple", "oracle", "enron inc.", "mcdonalds corporation")
y_ornames <- c("apple corp", "oracle inc", "enron", "mcdonalds co")
x <- data.frame("ornames_x" = x_ornames)
y <- data.frame("ornames_y" = y_ornames)

# Perform merge with fuzzy matching
linkedOrgs <- LinkOrgs(x = x,
                        y = y,
                        by.x = "ornames_x",
                        by.y = "ornames_y",
                        algorithm = "fuzzy",
                        MaxDist = 0.6)

print(linkedOrgs)
```

pDistMatch\_discrete     *Compute Discrete String Distances (Internal)*

## Description

Computes pairwise string distances between organization names in two data frames using discrete distance measures (e.g., Jaccard, OSA, Jaro-Winkler). Uses trigram indexing for efficient candidate filtering and parallel processing for speed.

## Usage

```
pDistMatch_discrete(
  x,
  y,
  by = NULL,
  by.x = NULL,
  by.y = NULL,
  embedDistMetric = NULL,
  return_stringdist = T,
  onlyUFT = T,
  qgram = 2,
  DistanceMeasure = "jaccard",
  MaxDist = 0.2,
  ReturnProgress = T,
  nCores = NULL,
  ReturnMaxDistThreshold = F
)
```

## Arguments

|      |  |
|------|--|
| x, y | Data frames containing organization names to be matched.   |
| by   | Character string; column name for matching when both data frames use the same column name. Overridden by by.x and by.y if specified. |

|                               |   |
|-------------------------------|---|
| <b>by.x</b>                   | Character string; column name in <b>x</b> containing organization names.  |
| <b>by.y</b>                   | Character string; column name in <b>y</b> containing organization names.  |
| <b>embedDistMetric</b>        | Not used in discrete matching (included for API consistency).   |
| <b>return_stringdist</b>      | Logical; if TRUE, returns string distances. Default is TRUE.  |
| <b>onlyUFT</b>                | Logical; if TRUE, processes only UTF-8 strings. Default is TRUE.  |
| <b>qgram</b>                  | Integer; the q-gram size for string distance calculation. Default is 2.   |
| <b>DistanceMeasure</b>        | Character; algorithm for computing pairwise string distances. Options include "jaccard", "osa", "jw". See <code>?stringdist::stringdist</code> for all options. Default is "jaccard". |
| <b>MaxDist</b>                | Numeric; maximum allowed distance between matched strings. Pairs with distances greater than this threshold are excluded. Default is 0.20.  |
| <b>ReturnProgress</b>         | Logical; if TRUE, progress information is available (currently disabled). Default is TRUE.  |
| <b>nCores</b>                 | Integer; number of CPU cores for parallel processing. Default is NULL (uses single core).   |
| <b>ReturnMaxDistThreshold</b> | Logical; if TRUE, returns the distance threshold used. Default is FALSE.  |

## Details

This function uses a two-stage approach for efficient matching:

1. **Trigram indexing:** Builds an index of character trigrams for each name, then filters candidate pairs to those sharing at least 5% of trigrams.
2. **Distance computation:** Computes exact string distances only for filtered candidates, returning pairs with distances at or below **MaxDist**.

The function automatically swaps **x** and **y** if **y** has fewer rows than **x** for more efficient parallelization.

## Value

A data frame with three columns:

- ix** Integer; row index in **x** of the matched record.
- iy** Integer; row index in **y** of the matched record.
- stringdist** Numeric; the string distance between the matched pair.

Returns an empty data frame if no matches are found below **MaxDist**.

## See Also

[pFuzzyMatch\\_discrete\(\)](#) for the higher-level wrapper that returns merged data, [stringdist::stringdist\(\)](#) for available distance measures.

## Examples

```
# Create synthetic data
x_orgnames <- c("apple", "oracle", "enron inc.", "mcdonalds corporation")
y_orgnames <- c("apple corp", "oracle inc", "enron", "mcdonalds co")
x <- data.frame("orgnames_x" = x_orgnames)
y <- data.frame("orgnames_y" = y_orgnames)

# Compute distances
distances <- pDistMatch_discrete(x = x,
                                    y = y,
                                    by.x = "orgnames_x",
                                    by.y = "orgnames_y",
                                    MaxDist = 0.5)
```

pDistMatch\_euclidean *Compute Euclidean Distances Between Embeddings (Internal)*

## Description

Computes pairwise Euclidean distances between embedding vectors from two sets of observations. This function is used internally for ML-based matching where organization names have been converted to numeric embeddings.

## Usage

```
pDistMatch_euclidean(
  embedx,
  embedy,
  MaxDist = NULL,
  embedDistMetric = NULL,
  ReturnProgress = T
)
```

## Arguments

|                 |  |
|-----------------|--|
| embedx          | Numeric matrix; embeddings for the first set of observations. Rows correspond to observations and columns to embedding dimensions.   |
| embedy          | Numeric matrix; embeddings for the second set of observations. Rows correspond to observations and columns to embedding dimensions.  |
| MaxDist         | Numeric; maximum allowed Euclidean distance. Pairs with distances greater than this threshold are excluded. If NULL, all pairs are returned.   |
| embedDistMetric | Optional function; custom distance metric. If NULL, Euclidean distance is computed. The function should take two arguments (expanded x vector and transposed y matrix) and return a distance vector. |
| ReturnProgress  | Logical; if TRUE, progress information is available (currently disabled). Default is TRUE.   |

## Details

This function computes Euclidean distances between all pairs of embedding vectors. For efficiency:

- Automatically swaps `embedx` and `embedy` if `embedy` has fewer rows for better vectorization.
- Uses JAX for GPU acceleration when available (detected automatically).
- Rounds embeddings to reduce precision overhead when embedding precision differs.

The function is typically called by [pFuzzyMatch\\_euclidean\(\)](#) rather than directly by users.

## Value

A data frame with three columns:

**ix** Integer; row index in `embedx` of the matched record.

**iy** Integer; row index in `embedy` of the matched record.

**stringdist** Numeric; the Euclidean distance between the matched pair's embeddings (named `stringdist` for consistency with discrete matching).

Returns an empty data frame if no matches are found below `MaxDist`.

## See Also

[pFuzzyMatch\\_euclidean\(\)](#) for the higher-level wrapper that returns merged data, [pDistMatch\\_discrete\(\)](#) for string-distance-based matching.

## Examples

```
## Not run:
# Create synthetic embeddings
embedx <- matrix(rnorm(4 * 256), nrow = 4)
embedy <- matrix(rnorm(4 * 256), nrow = 4)

# Compute distances
distances <- pDistMatch_euclidean(embedx = embedx,
                                    embedy = embedy,
                                    MaxDist = 5.0)

## End(Not run)
```

## Description

Performs fuzzy matching between two data frames using string distance measures (e.g., Jaccard, OSA, Jaro-Winkler). This is a wrapper around [pDistMatch\\_discrete\(\)](#) that returns the merged data frame with matched records.

**Usage**

```
pFuzzyMatch_discrete(
  x = NULL,
  by.x = NULL,
  embedx = NULL,
  y = NULL,
  by.y = NULL,
  embedy = NULL,
  embedDistMetric = NULL,
  MaxDist = NULL,
  qgram = 2,
  DistanceMeasure = "jaccard",
  AveMatchNumberPerAlias = NULL,
  nCores = NULL,
  ...
)
```

**Arguments**

|                        |   |
|------------------------|---|
| x, y                   | Data frames to be merged.   |
| by.x                   | Character string; column name in x containing organization names.   |
| embedx, embedy         | Optional embedding matrices (not used in discrete matching, included for API consistency).  |
| by.y                   | Character string; column name in y containing organization names.   |
| embedDistMetric        | Optional custom distance metric (not used in discrete matching).  |
| MaxDist                | Numeric; maximum allowed distance between matched strings. Pairs with distances greater than this threshold are excluded. If AveMatchNumberPerAlias is specified, it takes priority.  |
| qgram                  | Integer; the q-gram size for string distance calculation. Default is 2.   |
| DistanceMeasure        | Character; algorithm for computing pairwise string distances. Options include "jaccard", "osa", "jw". See <code>?stringdist::stringdist</code> for all options. Default is "jaccard". |
| AveMatchNumberPerAlias | Numeric; target average number of matches per alias. If specified, automatically calibrates MaxDist using <a href="#">GetCalibratedDistThres()</a> .                                  |
| nCores                 | Integer; number of CPU cores for parallel processing. Default is NULL (uses single core).   |
| ...                    | Additional arguments (currently unused).  |

**Details**

This function uses trigram indexing to efficiently filter candidate matches before computing exact string distances. This approach significantly speeds up matching for large datasets.

**Value**

A data frame containing matched records from x and y, with columns from both data frames (suffixed with .x and .y respectively) and a `stringdist` column indicating the distance between each matched pair.

**See Also**

[pDistMatch\\_discrete\(\)](#) for the underlying distance computation, [GetCalibratedDistThres\(\)](#) for automatic threshold calibration, [pFuzzyMatch\\_euclidean\(\)](#) for embedding-based matching.

**Examples**

```
# Create synthetic data
x_orgnames <- c("apple", "oracle", "enron inc.", "mcdonalds corporation")
y_orgnames <- c("apple corp", "oracle inc", "enron", "mcdonalds co")
x <- data.frame("orgnames_x" = x_orgnames)
y <- data.frame("orgnames_y" = y_orgnames)

# Perform fuzzy matching
matched <- pFuzzyMatch_discrete(x = x,
                                  y = y,
                                  by.x = "orgnames_x",
                                  by.y = "orgnames_y",
                                  MaxDist = 0.5)
```

**pFuzzyMatch\_euclidean** *Fuzzy Match with Euclidean Distance on Embeddings*

**Description**

Performs fuzzy matching between two data frames using Euclidean distance on pre-computed embeddings. This is a wrapper around [pDistMatch\\_euclidean\(\)](#) that returns the merged data frame with matched records.

**Usage**

```
pFuzzyMatch_euclidean(
  x = NULL,
  by.x = NULL,
  embedx = NULL,
  y = NULL,
  by.y = NULL,
  embedy = NULL,
  embedDistMetric = NULL,
  MaxDist = NULL,
  AveMatchNumberPerAlias = NULL,
  ...
)
```

**Arguments**

|                |   |
|----------------|---|
| x, y           | Data frames to be merged.   |
| by.x           | Character string; column name in x containing organization names.   |
| embedx, embedy | Numeric matrices containing embeddings for records in x and y respectively. Rows correspond to observations and columns to embedding dimensions. These are typically produced by the ML backend (see <a href="#">LinkOrgs()</a> with algorithm = "ml"). |

|                        |  |
|------------------------|--|
| by.y                   | Character string; column name in y containing organization names.  |
| embedDistMetric        | Optional custom distance metric function for computing distances between embeddings. If NULL, Euclidean distance is used.  |
| MaxDist                | Numeric; maximum allowed Euclidean distance between matched embeddings.<br>Pairs with distances greater than this threshold are excluded. If AveMatchNumberPerAlias is specified, it takes priority. |
| AveMatchNumberPerAlias | Numeric; target average number of matches per alias. If specified, automatically calibrates MaxDist using <a href="#">GetCalibratedDistThres()</a> .   |
| ...                    | Additional arguments (currently unused).   |

## Details

This function is typically used internally by `LinkOrgs()` when `algorithm = "ml"` or `DistanceMeasure = "ml"`. It computes Euclidean distances between embedding vectors rather than string distances.

## Value

A data frame containing matched records from x and y, with columns from both data frames (suffixed with .x and .y respectively) and a stringdist column indicating the Euclidean distance between each matched pair's embeddings.

#### **See Also**

`pDistMatch_euclidean()` for the underlying distance computation, `GetCalibratedDistThres()` for automatic threshold calibration, `pFuzzyMatch_discrete()` for string-distance-based matching.

## Examples

|                     |               |
|---------------------|---------------|
| <code>print2</code> | <i>print2</i> |
|---------------------|---------------|

## Description

Prints a message with a timestamp prefix.

## Usage

```
print2(text, quiet = F)
```

## Arguments

|                    |  |
|--------------------|--|
| <code>text</code>  | Character string to print.                           |
| <code>quiet</code> | Logical; if TRUE, suppress output. Default is FALSE. |

## Value

Invisibly returns NULL. Called for its side effect of printing.

## Examples

```
print2("Hello world!")
```

|                     |  |
|---------------------|--|
| <code>url2dt</code> | <i>Download CSV from URL to data.table</i> |
|---------------------|--|

## Description

Downloads a zipped CSV file from a URL and loads it into memory as a data.table. Automatically handles Dropbox URLs by converting them to direct download links.

## Usage

```
url2dt(url)
```

## Arguments

|                  |  |
|------------------|--|
| <code>url</code> | Character string; the URL pointing to a .csv.zip or .csv.gz file. Dropbox share links are automatically converted to direct download URLs. |
|------------------|--|

## Details

This function:

1. Converts Dropbox share links to direct download URLs using [dropboxURL2downloadURL\(\)](#)
2. Downloads the file to a temporary directory
3. Unzips (if .zip) or decompresses (if .gz) the file
4. Reads the CSV file using `data.table:::fread()`
5. Cleans up the temporary file

**Value**

A data.table containing the downloaded data.

**See Also**

[dropboxURL2downloadURL\(\)](#) for URL conversion, [data.table::fread\(\)](#) for the underlying CSV reader.

**Examples**

```
## Not run:  
# Download from Dropbox  
my_dt <- url2dt("https://www.dropbox.com/s/example/data.csv.zip?dl=0")  
  
## End(Not run)
```

# Index

AssessMatchPerformance, 2  
BuildBackend, 3  
data.table::fread(), 16, 17  
dropboxURL2downloadURL, 4  
dropboxURL2downloadURL(), 16, 17  
GetCalibratedDistThres, 5  
GetCalibratedDistThres(), 13–15  
LinkOrgs, 6  
LinkOrgs(), 4, 14, 15  
pDistMatch\_discrete, 9  
pDistMatch\_discrete(), 12, 14  
pDistMatch\_euclidean, 11  
pDistMatch\_euclidean(), 14, 15  
pFuzzyMatch\_discrete, 12  
pFuzzyMatch\_discrete(), 10, 15  
pFuzzyMatch\_euclidean, 14  
pFuzzyMatch\_euclidean(), 12, 14  
print2, 16  
stringdist::stringdist(), 10  
url2dt, 16  
url2dt(), 5